

CIS 520, Machine Learning, Fall 2015: Assignment 4  
Due: Thursday, October 15th, 11:59pm, PDF to Canvas

Arpit Panwar

## Boosting and MDL

### 1. Analyzing the training error of boosting

Consider the AdaBoost algorithm you saw in class. In this question we will try to analyze the training error of boosting.

1. Given a set of  $m$  examples,  $(x_i, y_i)$  ( $y_i$  is the class label of  $x_i$ ),  $i = 1, \dots, m$ , let  $h_t(x)$  be the weak classifier obtained at step  $t$ , and let  $\alpha_t$  be its weight. Recall that the final classifier is

$$H(x) = \text{sign}(f(x)), \text{ where } f(x) = \sum_{t=1}^T \alpha_t h_t(x).$$

Show that the training error of the final classifier can be bounded from above by an exponential loss function:

$$\frac{1}{m} \sum_{i=1}^m I(H(x_i) \neq y_i) \leq \frac{1}{m} \sum_{i=1}^m \exp(-f(x_i)y_i),$$

where  $I(a = b)$  is the indicator function. It is equal to 1 if  $a = b$ , and 0 otherwise

*Hint:*  $e^{-x} \geq 1 \Leftrightarrow x \leq 0$ .

$$\begin{aligned} D_{t+1}(i) &= \frac{D_{t-1}(i) * e^{-\alpha_t * y_i * h_t(x_i)} * e^{-\alpha_{t-1} * y_i * h_{t-1}(x_i)}}{Z_t * Z_{t-1}} \\ &\quad \text{Expanding till } D_1 \\ &= \frac{D_1(i) * e^{-\alpha_t * y_i * h_t(x_i)} * e^{-\alpha_{t-1} * y_i * h_{t-1}(x_i)} \dots e^{-\alpha_1 * y_i * h_1(x_i)}}{Z_t * Z_{t-1} * Z_{t-2} \dots * Z_1} \\ &= \frac{e^{\sum_t y_i * \alpha_t * h_t(x_i)}}{m * \prod_t Z_t} \\ D_{t+1}(i) &= \frac{e^{-y_i f(x_i)}}{m * \prod_t Z_t} \\ &\quad \text{When } H(x_i) \neq y_i \text{ } y_i * f(x_i) \leq 0. \text{ Thus } e^{-y_i * f(x_i)} \geq 1 \\ &\quad I(H(x_i) \neq y_i) \leq e^{-y_i * f(x_i)} \\ &\quad \frac{1}{m} * \sum_i I(H(x_i) \neq y_i) = \frac{1}{m} * \sum_i e^{-y_i * f(x_i)} \end{aligned}$$

2. Remember that

$$D_{t+1}(i) = \frac{D_t(i) \exp(-\alpha_t y_i h_t(x_i))}{Z_t}$$

Use this recursive definition to prove the following.

$$\frac{1}{m} \sum_{i=1}^m \exp(-f(x_i)y_i) = \prod_{t=1}^T Z_t, \quad (1)$$

where  $Z_t$  is the normalization factor for distribution  $D_{t+1}$ :

$$Z_t = \sum_{i=1}^m D_t(i) \exp(-\alpha_t y_i h_t(x_i)). \quad (2)$$

*Hint:* Remember that  $e^{\sum_i g_i} = \prod_i e^{g_i}$ ,  $D_1(i) = \frac{1}{m}$ , and that  $\sum_i D_{t+1}(i) = 1$ .

$$\begin{aligned} & \text{Expanding the above equation for } D_{t+1} \\ D_{t+1}(i) &= \frac{D_{t-1}(i) * e^{-\alpha_t * y_i * h_t(x_i)} * e^{-\alpha_{t-1} * y_i * h_{t-1}(x_i)}}{Z_t * Z_{t-1}} \\ & \text{Expanding till } D_1 \\ &= \frac{D_1(i) * e^{-\alpha_t * y_i * h_t(x_i)} * e^{-\alpha_{t-1} * y_i * h_{t-1}(x_i)} \dots e^{-\alpha_1 * y_i * h_1(x_i)}}{Z_t * Z_{t-1} * Z_{t-2} \dots * Z_1} \\ &= \frac{e^{\sum_t y_t * \alpha_t * h_t(x_i)}}{m * \prod_t Z_t} \\ &= \frac{e^{-y_i \sum_t \alpha_t h_t(x_i)}}{m * \prod_t Z_t} \\ & \text{We know that } f(x) = \sum_t \alpha_t * h_t(x_i) \\ D_{t+1}(i) &= \frac{e^{-y_i f(x_i)}}{m * \prod_t Z_t} \\ & \text{Summing over all } i \\ \sum_i D_{t+1}(i) &= \frac{\sum_i e^{-y_i f(x_i)}}{m * \prod_t Z_t} \\ & \text{We know } \sum_i D_{t+1}(i) = 1. \text{ Substituting and rearranging} \\ \prod_t Z_t &= \sum_i e^{-y_i f(x_i)} \end{aligned}$$

3.

$$\frac{1}{m} \sum_{i=1}^m \exp(-f(x_i)y_i) = \prod_{t=1}^T Z_t, \quad (3)$$

suggests that the training error can be reduced rapidly by greedily optimizing  $Z_t$  at each step. You have shown that the error is bounded from above:

$$\epsilon_{\text{training}} \leq \prod_{t=1}^T Z_t.$$

Observe that  $Z_1, \dots, Z_{t-1}$  are determined by the first  $(t-1)$  rounds of boosting, and we cannot change them on round  $t$ . A greedy step we can take to minimize the training error bound on round  $t$  is to minimize  $Z_t$ .

In this question, you will prove that for binary weak classifiers,  $Z_t$  from

$$Z_t = \sum_{i=1}^m D_t(i) \exp(-\alpha_t y_i h_t(x_i)). \quad (4)$$

is minimized by picking  $\alpha_t$  as:

$$\alpha_t^* = \frac{1}{2} \log \left( \frac{1 - \epsilon_t}{\epsilon_t} \right), \quad (5)$$

where  $\epsilon_t$  is the training error of weak classifier  $h_t$  for the weighted dataset:

$$\epsilon_t = \sum_{i=1}^m D_t(i) I(h_t(x_i) \neq y_i).$$

where  $I$  is the indicator function. For this proof, only consider the simplest case of binary classifiers, i.e. the output of  $h_t(x)$  is binary,  $\{-1, +1\}$ .

For this special class of classifiers, first show that the normalizer  $Z_t$  can be written as:

$$Z_t = (1 - \epsilon_t) \exp(-\alpha_t) + \epsilon_t \exp(\alpha_t).$$

*Hint:* Consider the sums over correctly and incorrectly classified examples separately.

$$Z_t = \sum_{i=1}^m D_t(i) \exp(-\alpha_t y_i h_t(x_i))$$

*Splitting into with  $h(x_i) \neq y_i$  and  $h(x_i) = y_i$*

$$Z_t = \sum_{h(x_i)=y_i} D_t(x_i) e^{-y_i * \alpha_t * h_t(x_i)} + \sum_{h(x_i) \neq y_i} D_t(x_i) e^{-y_i * \alpha_t * h_t(x_i)}$$

$$\text{We know that } \epsilon_t = \sum_{i=1}^m D_t(i) I(h_t(x_i) \neq y_i).$$

*Additionally we know  $e^{-y_i * \alpha_t * h_t(x_i)} = e^{-\alpha_t}$  when  $h_t(x_i) = y_i$  and  $e^{-y_i * \alpha_t * h_t(x_i)} =$*

*Also since the classifier is binary we when we classify correctly the error is  $(1 - \epsilon_t)$*

$$Z_t = (1 - \epsilon_t) e^{-\alpha_t} + \epsilon_t e^{\alpha_t}$$

4. Now, prove that the value of  $\alpha_t$  that minimizes this definition of  $Z_t$  is given by Equation

$$\alpha_t^* = \frac{1}{2} \log \left( \frac{1 - \epsilon_t}{\epsilon_t} \right). \quad (6)$$

*Using the equation above and taking differential w.r.t  $\alpha_t$  and setting to 0*

$$\frac{dZ_t}{d\alpha_t} = \frac{d}{d\alpha_t} (1 - \epsilon_t) e^{-\alpha_t} + \frac{d}{d\alpha_t} \epsilon_t * e^{-\alpha_t}$$

$$0 = - (1 - \epsilon_t) e^{-\alpha_t} + \epsilon_t e^{\alpha_t}$$

$$\epsilon_t e^{\alpha_t} = (1 - \epsilon_t) e^{-\alpha_t}$$

$$e^{\alpha_t} = \frac{(1 - \epsilon_t) * e^{-\alpha_t}}{\epsilon_t}$$

$$e^{2\alpha_t} = \frac{1 - \epsilon_t}{\epsilon_t}$$

$$2\alpha_t = \log \frac{1 - \epsilon_t}{\epsilon_t}$$

$$\alpha_t = \frac{1}{2} * \log \frac{1 - \epsilon_t}{\epsilon_t}$$

5. Prove that for the above value of  $\alpha_t$  we have  $Z_t = 2\sqrt{\epsilon_t(1 - \epsilon_t)}$ .

Substituting the values from 1.3 into 1.4

$$\begin{aligned}
 Z_t &= \frac{1 - \epsilon_t}{\sqrt{\frac{1 - \epsilon_t}{\epsilon_t}}} + \epsilon_t * \sqrt{\frac{1 - \epsilon_t}{\epsilon_t}} \\
 &= \frac{1 - \epsilon_t}{\sqrt{\frac{1 - \epsilon_t}{\epsilon_t}}} + \sqrt{\epsilon_t * (1 - \epsilon_t)} \\
 &= \sqrt{\epsilon_t * (1 - \epsilon_t)} + \sqrt{\epsilon_t * (1 - \epsilon_t)} \\
 &= 2 * \sqrt{\epsilon_t * (1 - \epsilon_t)}
 \end{aligned}$$

6. Furthermore, let  $\epsilon_t = \frac{1}{2} - \gamma_t$  and prove that  $Z_t \leq \exp(-2\gamma_t^2)$ . Therefore, we will know

$$\epsilon_{training} \leq \prod_t Z_t \leq \exp(-2 \sum_t \gamma_t^2).$$

*Hint:*  $\log(1 - x) \leq -x$  for  $0 < x \leq 1$ .

Using 1.5 and substituting the value for  $\epsilon_t$

$$\begin{aligned}
 Z_t &= 2\sqrt{(\frac{1}{2} - \gamma_t) * (\frac{1}{2} + \gamma_t)} \\
 &= 2\sqrt{\frac{1}{4} - \gamma_t^2} \\
 &= \sqrt{1 - 4\gamma_t^2} \\
 &= e^{\frac{1}{2} \log 1 - 4\gamma_t^2} \\
 &\quad \text{Using } \log(1 - x) \leq -x \\
 Z_t &\leq e^{\frac{1}{2} * 4 * -\gamma_t^2} \\
 Z_t &\leq e^{-2\gamma_t^2} \\
 \prod_t Z_t &\leq e^{-\sum_t 2\gamma_t^2}
 \end{aligned}$$

7. For every step of AdaBoost in a binary classifier if we have a classifier with  $\epsilon_t > 0.5$  then there will always be a classifier which classifies exactly opposite and has an error  $1 - \epsilon_t$ .

Additionally if  $\epsilon_t > 0.5$  we get  $\alpha_t < 0$  which will penalize the correct prediction and not the once predicted incorrectly. And since we know that we always converge the weak classifier should have error  $< 0.5$ .

8. For  $\epsilon_t = 0.5$ , using the equation for  $\alpha_t$  above, we get  $\alpha_t = 0$ . Thus  $D_{t+1} = D_t$  thus the error gets stuck.

## 2. Adaboost on a toy dataset

Now we will apply Adaboost to classify a toy dataset. Consider the dataset shown in Figure 1a. The dataset consists of 4 points:  $(X_1 : 0, -1, -)$ ,  $(X_2 : 1, 0, +)$ ,  $(X_3 : -1, 0, +)$  and  $(X_4 : 0, 1, -)$ . For this part, you may find it helpful to use MATLAB as a calculator rather than doing the computations by hand. (You do not need to submit any code though.)




Figure 1: a) Toy data in Question . b)  $h_1$  in Question




Figure 2: a) H1 b) H2 c) H3 d) H4

1. For  $T = 4$ , show how Adaboost works for this dataset, using simple decision stumps (depth-1 decision trees that simply split on a single variable once) as weak classifiers. For each timestep compute the following:

$$\epsilon_t, \alpha_t, Z_t, D_t(i) \forall i,$$

Also for each timestep draw your weak classifier. For example  $h_1$  can be as shown in 1b).

	T=1	T=2	T=3	T=4
E	0.2500	0.1667	0.1000	0.0556
Alpha	0.5493	0.8047	1.0986	1.4166
D	[0.25;0.25;0.25;0.25]	[0.1667 ; 0.1667 ; 0.5000 ; 0.1667]	[0.1 ; 0.1 ; 0.3 ; 0.5]	[0.0556 ; 0.500 ; 0.1667 ; 0.2778]
Z	0.8660	0.7454	0.6000	0.4581

Table 1: Errors and Alphas

2. Training error is 0
3. The dataset is not linearly separable. Since the dataset is not linearly separable, decision tree can not make a better decision on non-linear dataset.

### 3. MDL on a toy dataset

Consider the following problem

We provide a data set generated from a particular model with  $N = 64$ , where we want to estimate

$$\hat{y} = w_1x_1 + w_2x_2 + w_3x_3$$

We want to use MDL to find the 'optimal'  $L_0$ -penalized model.

1. Estimate the three linear regressions

(We could actually try all possible subsets here, but instead we'll just try three.)

$$y_1 = w_1x_1$$

$$y_2 = w_1x_1 + w_2x_2$$

$$y_3 = w_1x_1 + w_2x_2 + w_3x_3$$

For each of the three cases, what is

- (a) the sum of square error
    - i)  $\text{Err}_1 = 1.2779e + 03$
    - ii)  $\text{Err}_2 = 835.0558$
    - iii)  $\text{Err}_3 = 834.7418$
  - (b) 2 times the estimated bits to code the residual ( $n \log \frac{\text{Error}}{n}$ )
    - i)  $\text{ERR\_bits}_1 = 191.6237$
    - ii)  $\text{ERR\_bits}_2 = 164.3914$
    - iii)  $\text{ERR\_bits}_3 = 164.3973$
  - (c) 2 times the estimated bits to code each residual plus model under AIC ( $2 * 1$  bit to code each feature)
    - i)  $\text{AIC\_bits}_1 = 193.6237$
    - ii)  $\text{AIC\_bits}_2 = 168.3914$
    - iii)  $\text{AIC\_bits}_3 = 170.3673$
  - (d) 2 times the estimated bits to code each residual plus model under BIC ( $2 * (1/2)\log(n)$  bits to code each feature)
    - i)  $\text{BIC\_bits}_1 = 195.7826$
    - ii)  $\text{BIC\_bits}_2 = 172.7092$
    - iii)  $\text{BIC\_bits}_3 = 176.8440$
2. Which model has the smallest minimum description length?
    - a) for AIC :- Model 2 gives minimum MDL
    - b) for BIC :- Model 2 gives minimum MDL
  3.
    - i. Error for model 1 =  $1.7786e + 03$
    - ii. Error for model 2 =  $1.1671e + 03$
    - iii. Error for model 3 =  $1.1727e + 03$